

## Structured Query Language:

### ● 1. What is SQL?

- SQL is a language used to communicate with relational databases. It allows you to create, read, update, and delete data — commonly known as CRUD operations.
- SQL or Relational databases are used to store and manage the data objects that are related to one another.
- A system used to manage these **relational databases** is known as Relational Database Management System (RDBMS).

### ● What Is a Database?

- A database is a structured collection of data that is stored in a computer system.
- A Database may contain many tables.

### ● What is RDBMS?

- RDBMS stands for Relational Database Management System. RDBMS is the basis for SQL, and for all modern database systems like MS SQL Server, IBM DB2, Oracle, MySQL, and Microsoft Access.

### ● What is a Table?

- The data in an RDBMS is stored in database objects known as tables. This table is basically a collection of related data entries and it consists of numerous columns and rows.
- a table is the most common and simplest form of data storage in a relational database

| ID | Name    | Age | Salary  | City       | Country |
|----|---------|-----|---------|------------|---------|
| 1  | Ramesh  | 32  | 2000.00 | Hyderabad  | India   |
| 2  | Mukesh  | 40  | 5000.00 | New York   | USA     |
| 3  | Sumit   | 45  | 4500.00 | Muscat     | Oman    |
| 4  | Kaushik | 25  | 2500.00 | Kolkata    | India   |
| 5  | Hardik  | 29  | 3500.00 | Bhopal     | India   |
| 6  | Komal   | 38  | 3500.00 | Saharanpur | India   |
| 7  | Ayush   | 25  | 3500.00 | Delhi      | India   |
| 8  | Javed   | 29  | 3700.00 | Delhi      | India   |

- What is a Field → like ID, Name, Age, Salary, City and Country.
- What is a Record or a Row?
  - A record is also called as a row of data is each individual entry that exists in a table

| ID | Name   | Age | Salary  | City      | Country |
|----|--------|-----|---------|-----------|---------|
| 1  | Ramesh | 32  | 2000.00 | Hyderabad | India   |

### ● Keys?

- **UNIQUE Key:** Ensures that all the values in a column are different. here null can also present.

```
CREATE TABLE table_name(
    column1 datatype UNIQUE KEY,
    column2 datatype UNIQUE KEY,
    .....
    .....
    columnN datatype
);
```

- **Primary Key:** A Primary Key is a column that uniquely identifies each row in a table.

- **Characteristics:**

- Must be **unique** for each row.
    - Cannot be **NULL**.
    - Each table can have **only one** primary key.

- CREATE TABLE Customers (
 CustomerID INT PRIMARY KEY,
 Name VARCHAR(100),
 Email VARCHAR(100)
 );

- **CustomerID** is the primary key — no two customers can have the same **CustomerID**.

- **Foreign Key:** A **Foreign Key** is a column that **references the primary key** of another table. It's used to create a **relationship between two tables**.

- **Characteristics:**


- Ensures **referential integrity** between related tables.
    - Can have **duplicate** values.
    - Can contain **NULLs**, unless specified otherwise.

- CREATE TABLE Orders (
 OrderID INT PRIMARY KEY,
 OrderDate DATE,
 CustomerID INT,
 FOREIGN KEY (CustomerID) REFERENCES
 Customers(CustomerID)
 );

In this example:

- **CustomerID** in **Orders** is a foreign key.
    - It references **CustomerID** in the **Customers** table.
    - This means each order must be linked to a valid customer.

- **Candidate Key** — any column or set of columns that can uniquely identify each row.
  - There can be **one or more candidate keys** in a table. But **only one** of them is chosen as the **Primary Key**.
  - So Candidate Key Satisfy all conditions of primary key.



- 
- **Composite Primary Key**
  - In many-to-many relationships (like students enrolled in courses), **no single column uniquely identifies a row**.
  - So we **combine two columns** to form a **composite primary key**.

Enrollment Table (Before Normalization)

| StudentID | CourseID | Grade |
|-----------|----------|-------|
| 1         | CS101    | A     |
| 2         | CS101    | B     |
| 1         | MATH201  | A-    |

Here:

- No single column (StudentID or CourseID) can uniquely identify a row.
  - But the pair (StudentID, CourseID) is unique for each row.
- ✓ So (StudentID, CourseID) is a **composite primary key**.

## ● Database Normalization:

- Database normalization is the process of efficiently organizing data in a database. There are two reasons of this normalization process –
  1. Eliminating redundant data, for example, storing the same data in more than one table.
  2. Ensuring data dependencies make sense.
- Both these reasons are worthy goals as they reduce the amount of space a database consumes and ensures that data is logically stored. Normalization consists of a series of guidelines that help guide you in creating a good database structure.
- Normalization guidelines are divided into normal forms; think of a form as the format or the way a database structure is laid out.
- **Normal Forms (NF):**
  - **First Normal Form (1NF)**
    - **Rule: No repeating groups or arrays. Each field should hold atomic (indivisible) values.**

✗ Not in 1NF:

| OrderID | CustomerName | Products      |
|---------|--------------|---------------|
| 1       | Alice        | Pen, Notebook |

✓ In 1NF:

| OrderID | CustomerName | Product  |
|---------|--------------|----------|
| 1       | Alice        | Pen      |
| 1       | Alice        | Notebook |

- **Second Normal Form (2NF)**

- A table is in 2NF if:
  1. It is already in 1NF.
  2. Every non-key column is fully dependent on the whole primary key, not just part of it.
  3. A **partial dependency** happens when a column depends **only on part** of a **composite primary key**, not the whole key.

🔑 Original Table (1NF but Not in 2NF)

| StudentID | CourseID | StudentName | CourseName | Instructor |
|-----------|----------|-------------|------------|------------|
| 1         | CS101    | Alice       | SQL Basics | Mr. Smith  |
| 2         | CS101    | Bob         | SQL Basics | Mr. Smith  |

- 4.
5. **Composite Primary Key: (StudentID, CourseID)**

**Problem:**

- a. **StudentName depends only on StudentID**
- b. **CourseName and Instructor depend only on CourseID.**
- c. ➡ So we have **partial dependencies** – columns depending on **part** of the key, not the **whole key**.

## ✅ Converting to 2NF (Remove Partial Dependencies)

### ◆ Step 1: Split the data into separate tables

#### 👤 Students Table

| StudentID | StudentName |
|-----------|-------------|
| 1         | Alice       |
| 2         | Bob         |

#### 📚 Courses Table

| CourseID | CourseName | Instructor |
|----------|------------|------------|
| CS101    | SQL Basics | Mr. Smith  |

#### 📋 Enrollment Table (links students to courses)

| StudentID | CourseID |
|-----------|----------|
| 1         | CS101    |
| 2         | CS101    |

6.
  - a. Now each non-key column depends fully on its primary key. → StudentName depends on StudentID , Instructor and CourseName depends on CourseID , CourseID depends on StudentID.
  - b. No partial dependencies remain.
- If the **primary key is just one column**, you're **already in 2NF** (because partial dependency is not possible).
  - But if the **key is composite**, like (StudentID, CourseID), then:
    - Every non-key column must depend on **both** parts.
    - If a column depends only on StudentID, it **violates 2NF**.

## • What is 3NF?

### ○ A table is in Third Normal Form (3NF) if:

1. It is in Second Normal Form (2NF) ✅
2. There is no transitive dependency — i.e., non-key columns do not depend on other non-key columns (or) non-key columns should depend on key columns only.

### • What is a transitive dependency?

- A transitive dependency means a non-key column depends on another non-key column, not directly on the primary key.

## Example of a Table **NOT** in 3NF

| EmpID | EmpName | DeptID | DeptName |
|-------|---------|--------|----------|
| 101   | Alice   | D01    | HR       |
| 102   | Bob     | D02    | IT       |

- **Primary Key:** EmpID
- EmpName and DeptID depend on EmpID ✓
- But DeptName depends on DeptID and DeptID depends on EmpID ✗
- EmpID → DeptID → DeptName (transitive dependency)
- So, this is Not in 3NF

### ✓ How to Fix It (Convert to 3NF)

Split the table into two:

#### 1. Employees table

| EmpID | EmpName | DeptID |
|-------|---------|--------|
| 101   | Alice   | D01    |
| 102   | Bob     | D02    |

#### 2. Departments table

| DeptID | DeptName |
|--------|----------|
| D01    | HR       |
| D02    | IT       |

Now:

- DeptName depends only on DeptID, not on EmpID
- No transitive dependencies ✓
- Table is now in 3NF

### ■ 2 NF VS 3 NF: Let Consider a question Which

- ✓ Satisfy 2NF: No partial dependencies
- ✗ Fail 3NF: Has a transitive dependency

#### Example Table: Employees

| EmpID | EmpName | DeptID | DeptName |
|-------|---------|--------|----------|
| 101   | Alice   | D01    | HR       |
| 102   | Bob     | D02    | IT       |
| 103   | Carol   | D01    | HR       |

- 🔑 Primary Key: EmpID (a single-column key)

- **Why This is in 2NF:**
  - All non-key attributes (*EmpName*, *DeptID*, *DeptName*) depend on the whole primary key (*EmpID*).
  - No partial dependency exists (since key is not composite).
  - *DeptName* depends on *DeptID* and *DeptID* depends on *EmpID* ,However it indirectly depends on *EmpID*.

- **Why This is NOT in 3NF:**
  - *DeptName* depends on *DeptID*, not directly on *EmpID*.
  - But *DeptID* depends on *EmpID* → So, there's a transitive dependency:

■ **How to Fix It (Convert to 3NF):**

- ◆ Split into two tables:

1. **Employees :**

| EmpID | EmpName | DeptID |
|-------|---------|--------|
| 101   | Alice   | D01    |
| 102   | Bob     | D02    |
| 103   | Carol   | D01    |

2. **Departments :**

| DeptID | DeptName |
|--------|----------|
| D01    | HR       |
| D02    | IT       |

- Now:
  - All non-key columns depend **only on Primary keys**
  - No transitive dependencies ✅ → **3NF achieved**

- **What is Denormalization in DBMS:**Denormalization is the process of intentionally introducing redundancy into a database by combining tables that were separated during normalization. This is done mainly to improve performance, especially read/query speed.

→ **Definition:**Denormalization is the process of joining normalized tables into fewer tables by adding redundant data, in order to reduce JOINS and improve query performance.

→ **Why Denormalize?**

While normalization helps reduce redundancy and improve data integrity, it can lead to:

- Many small tables
- Complex JOIN operations
- Slower read performance

→ **Denormalization improves:**


- Query speed (fewer JOINS)
- Read-heavy performance

→ **But it sacrifices:**

- Storage efficiency
- Data consistency
- Update performance (more risk of anomalies)

## Example:

### Normalized Tables:

 Normalized Tables:

#### 1. Customers

| CustomerID | Name  |
|------------|-------|
| 1          | Alice |
| 2          | Bob   |

#### 2. Orders

| OrderID | CustomerID | Product |
|---------|------------|---------|
| 101     | 1          | Laptop  |
| 102     | 2          | Phone   |

**To get the customer's name for an order, you'd need a JOIN:**

```
■ SELECT Orders.OrderID, Customers.Name, Orders.Product
FROM Orders
JOIN Customers ON Orders.CustomerID =
Customers.CustomerID;
```

### Denormalized Table:



| OrderID | CustomerID | Name  | Product |
|---------|------------|-------|---------|
| 101     | 1          | Alice | Laptop  |
| 102     | 2          | Bob   | Phone   |

- **No JOIN needed**
- **But now Customer Name is duplicated**
- **If Alice changes her name, it must be updated in every order row**

### Key Trade-offs:

| Benefit               | Cost                                |
|-----------------------|-------------------------------------|
| Faster SELECT queries | Higher storage (redundancy)         |
| Fewer JOINS           | More chance of <b>inconsistency</b> |
| Simpler queries       | Harder updates (data is repeated)   |

### When to Use Denormalization:

- In read-heavy systems (data warehouses, analytics)
- When JOINS slow down performance
- When data changes rarely
- For reporting and summary tables
- **What is a Transaction:**
  - A **transaction** is a sequence of one or more database operations (like insert, update, delete) executed as a **single logical unit of work**.
  - **Example of a Transaction:**

Imagine transferring money from Account A to Account B:

- Step 1: Deduct \$100 from Account A
- Step 2: Add \$100 to Account B

Both steps must succeed together. If step 2 fails, step 1 should be undone to keep data consistent.

### Why Use Transactions?

- To **maintain data integrity** during complex operations
- To **handle failures gracefully** (rollback on errors)
- To **enable concurrent access safely**

- **Transaction States**

- A transaction goes through several states during its lifecycle:

| State               | Description   |
|---------------------|---|
| Active              | Transaction is executing its operations.                  |
| Partially Committed | All operations are done, waiting to commit.               |
| Committed           | Changes are permanently saved in the database.            |
| Failed              | Transaction encountered an error and must be rolled back. |

- Aborted Transaction rolled back, database restored to previous state.

- **Concurrency Control**

When multiple transactions run **simultaneously**, concurrency control ensures **correctness and consistency** by handling:

- **Isolation:** Transactions should not interfere with each other.
- **Serializability:** The effect of concurrent transactions is the same as if they ran one after the other.

(Or) Handling **concurrency** in a DBMS is crucial to ensure **data consistency, correctness, and isolation** when **multiple transactions access the database at the same time**.

- **What Is Concurrency Control?**

**Concurrency control** ensures that when multiple transactions run concurrently, the result is **correct** and **the same as if they were run one after the other (serially)**.

⚠ **Problems Without Concurrency Control:**

| Problem             | Description  |
|---------------------|--|
| Dirty Read          | A transaction reads data modified by another transaction that hasn't committed.  |
| Lost Update         | Two transactions read the same data and both update it, but one update is lost.  |
| Non-repeatable Read | A transaction reads the same row twice and gets different data because another transaction modified it in between.                   |
| Phantom Read        | A transaction reads a set of rows matching a condition, but another transaction inserts new rows that now also match that condition. |

**Common techniques:**

**1. Lock-Based Protocols**

- Locks prevent other transactions from accessing data being used.

| Lock Type          | Description                              |
|--------------------|--|
| Shared Lock (S)    | Allows reading; multiple can coexist.    |
| Exclusive Lock (X) | Allows read and write; only one allowed. |

## Two-Phase Locking (2PL)

- Phase 1: Growing – acquire all needed locks.
- Phase 2: Shrinking – release locks; no new locks can be acquired.  
This ensures serializability.

### Transaction T1:

```
sql

START TRANSACTION;
SELECT balance FROM accounts WHERE id = 1 FOR UPDATE;
-- Lock is acquired here!
UPDATE accounts SET balance = balance - 100 WHERE id = 1;
COMMIT;
```

### Transaction T2 (starts before T1 commits):

```
sql

START TRANSACTION;
SELECT balance FROM accounts WHERE id = 1 FOR UPDATE;
-- This will wait (or be blocked) until T1 releases the lock
UPDATE accounts SET balance = balance + 50 WHERE id = 1;
COMMIT;
```

- What happens:
  - T2 waits for T1 to release the exclusive lock.
  - No dirty read or lost update occurs.
- In SQL, locks are implicitly acquired when using certain clauses like:

♦ **SELECT ... FOR UPDATE**

This tells the DBMS:

“I want to read this row and I plan to update it, so lock it exclusively.”

♦ **UPDATE, INSERT, or DELETE**

These automatically acquire exclusive locks on affected rows or tables

- **Timestamp-Based Protocol?**

Each transaction is assigned a unique timestamp when it starts. This timestamp determines the order in which transactions should logically execute — older transactions go first. It ensures serializability without using locks.

## Key Rules:

Each data item (e.g., variable or row) tracks:

- **read\_TS(X): Latest timestamp of any transaction that read X.**
- **write\_TS(X): Latest timestamp of any transaction that wrote to X.**

Then, for any transaction T:

| Operation             | Rule to check   |
|-----------------------|---|
| <code>read(X)</code>  | Allowed if T's timestamp $\geq$ <code>write_TS(X)</code>                                  |
| <code>write(X)</code> | Allowed if T's timestamp $\geq$ both <code>read_TS(X)</code> and <code>write_TS(X)</code> |

- If a rule is violated, the transaction is **rolled back (aborted)** and restarted with a **new timestamp**.

## Example:

We have a data item X (e.g., a bank balance) starting at ₹1000.

Initial timestamps:

- `read_TS(X) = 0`
- `write_TS(X) = 0`

## Transaction Start Times

| Transaction | Timestamp |
|-------------|-----------|
| T1          | 5         |
| T2          | 10        |
| T3          | 15        |

## Transaction Actions

| Step | Transaction | Operation | Timestamp | read_TS(X) | write_TS(X) | Action   | Result   |
|------|-------------|-----------|-----------|------------|-------------|--|----------|
| 1    | T1          | read(X)   | 5         | 0          | 0           | $5 \geq 0$ ✓                                   | Allowed  |
|      |             |           |           | → 5        |             | Update read_TS                                 |          |
| 2    | T2          | write(X)  | 10        | 5          | 0           | $10 \geq 5$ and $10 \geq 0$ ✓                  | Allowed  |
|      |             |           |           | 5          | → 10        | Update write_TS                                |          |
| 3    | T3          | read(X)   | 15        | 5          | 10          | $15 \geq 10$ ✓                                 | Allowed  |
|      |             |           |           | → 15       | 10          | Update read_TS                                 |          |
| 4    | T1          | write(X)  | 5         | 15         | 10          | $5 < 15$ ✗ (read_TS) and $5 < 10$ ✗ (write_TS) | Rejected |
|      |             |           |           |            |             | Rollback T1                                    | ✗        |
| 5    | T3          | write(X)  | 15        | 15         | 10          | $15 \geq 15$ and $15 \geq 10$ ✓                | Allowed  |
|      |             |           |           | ↓ 15       | → 15        | Update write_TS                                | ✓        |

### Final State of X:

- **read\_TS(X) = 15**
- **write\_TS(X) = 15**
- **Last successful write by T3**

### Why Use a Timestamp-Based Protocol?

- **No need for locks**
- **No waiting or deadlocks**
- **Ensures serializability**
- **Works well when conflicts are rare and reads dominate**

### What Is Optimistic Concurrency Control (OCC)?

Optimistic Concurrency Control assumes: Conflicts between transactions are rare, so don't lock anything upfront.

Instead of using locks, OCC allows transactions to execute freely and only check for conflicts at the end, during a validation phase.

### OCC Phases:

Each transaction goes through three phases:

#### 1. Read Phase

- The transaction reads data and performs computations.

- All updates are stored locally (not written to the database yet).

## 2. Validation Phase

- Before committing, the system checks:  
"Did any other transaction modify the data I used?"

## 3. Write Phase

- If validation succeeds, the updates are written to the database.
- If validation fails, the transaction is rolled back and restarted.

### Example:

- **Two Transactions: T1, T2,**
- Let's say we have a table with one row:

| ID | Balance |
|----|---------|
| 1  | ₹1000   |

### Scenario

#### ● T1 starts:

- **Read Phase:** Reads **balance = 1000** and stores it locally.
- Do some calculations (e.g., withdraw ₹100).

#### ● T2 starts during T1:

- **Read Phase:** Also reads **balance = 1000**.
- Deposits ₹200.

#### ● T2 ends and enters Validation Phase:

- No one has written anything yet, so validation passes ✓
- **Write Phase:** Balance becomes ₹1200.

#### ● T1 now tries to commit:

- **Validation Phase:** Compares **balance = 1000** (what T1 read) with current DB value (₹1200)
- **Mismatch detected ✗** — someone else (T2) modified the data after T1 read it.
- → T1 is aborted and restarted

## Why Use OCC?

- Best for systems with many reads and few writes (e.g., analytics dashboards).
- Avoids locking overhead.
- Prevents deadlocks.

### OCC vs. 2PL vs. Timestamp

| Feature        | OCC                     | 2PL (Locks)      | Timestamp             |
|----------------|-------------------------|------------------|-----------------------|
| Uses Locks     | ✗ No                    | ✓ Yes            | ✗ No                  |
| Deadlock       | ✗ Never                 | ✓ Possible       | ✗ Never               |
| Best Use Case  | Read-heavy systems      | Mixed workloads  | Ordered access needed |
| Conflict Check | At the end (validation) | During execution | At each read/write    |

## 4. Multiversion Concurrency Control (MVCC) [Same Like ReadWriteArrayList in Java concurrency]

MVCC allows multiple transactions to access the **same data at the same time** by storing **multiple versions** of a data item.

It solves the problem of **read-write conflicts** by giving:

- **Readers a consistent snapshot** of the data.
- **Writers a new version** of the data, without blocking readers.

You can **read old versions** of data without waiting for writes to finish. (Or)

### Idea:

- **Each write creates a new version.**
- **Readers see the data as it was when their transaction started.**

### Example:

Initial Table:

| id | balance | version |
|----|---------|---------|
| 1  | 1000    | 1       |

**T1 starts: Reads version 1**

**T2 starts after T1:**

- `UPDATE accounts SET balance = 1100 WHERE id = 1;`

**DBMS creates a new version (version 2) with balance = 1100  
T1 still sees version 1 (balance = 1000)**

- **Even though T2 updated it**

**This prevents non-repeatable reads and allows true read and write**

**isolation**

## ■ **Key Properties of Transactions (ACID):**

### 1. **Atomicity**

- All steps in a transaction are completed or none are.
- If any step fails, the entire transaction is rolled back.

#### **Example:**

Transferring \$100 from Account A to Account B.

- Step 1: Deduct \$100 from Account A
- Step 2: Add \$100 to Account B

If Step 2 fails after Step 1 succeeded, the database must undo Step 1 (rollback) so the money isn't lost.

### 2. **Consistency**

- A transaction takes the database from one **valid state** to another, maintaining all rules and constraints.

#### **Example:**

Suppose there's a rule: **Account balance cannot be negative.**

Before transaction:

Account A balance = \$200

Transaction:

Withdraw \$250

- This violates the rule → transaction is aborted or rolled back.(here the valid state is greater than or equal to 0).
- The database remains consistent (no negative balance).

### 3. **Isolation**

- Transactions run independently without interference from others.
- Intermediate states are not visible to other transactions.



- **Example:**

Two transactions run simultaneously:

- T1 transfers \$100 from Account A to B.
- T2 reads the balance of Account A.

Isolation ensures T2 sees either the balance **before** or **after** T1, but never an inconsistent intermediate state (like after deduction but before addition).

#### 4. **Durability**

- Once a transaction is committed, its changes are permanent, even if there is a system failure.

**Example:**

Once a transaction commits, like a money transfer, the changes persist even if:

- The system crashes immediately after commit.
- Power failure happens.

The DBMS ensures the changes are saved on disk and restored when back online.

- **What is SQL Syntax?**

- **Case Sensitivity**

- The most important point to be noted here is that SQL is case insensitive, which means **SELECT** and **Select** have same meaning in SQL statements
- But column names and values are case insensitive. Ex: EmpID(column Name) should be EmpID not empID or Empid or anyother.

- All the SQL statements require a **semicolon (;)** at the end of each statement.

- **SQL CREATE DATABASE Statement:**

- `CREATE DATABASE databasename;`

- **SQL USE Statement:**

- Once the database is created, it needs to be used in order to start storing the data accordingly.
- Database is collection tables (keep in mind).
- you can work with a database without explicitly using **USE**, but:

- In SQL tools like MySQL, SQL Server, or MariaDB:
  - `USE database_name;` tells the system **which database to run queries on.**
  - If you **don't use USE**, then:
    - You must **specify the database name** each time like this:
- `SELECT * FROM my_database.customers;`
- `My_database` → It is Database.
- `Customers` → It is table.

#### ◦ SQL Backup Database Statement

- `BACKUP DATABASE database_name  
TO DISK = 'filepath'  
GO`

#### ◦ Restore Database From Backup

- `RESTORE DATABASE database_name  
FROM DISK = 'filepath';  
GO`

#### ◦ SQL CREATE TABLE Statement:

- `CREATE TABLE table_name(  
    column1 datatype,  
    column2 datatype,  
    column3 datatype,  
    .....  
    columnN datatype,  
    PRIMARY KEY( one or more columns )  
);`

#### ■ SQL CREATE TABLE IF NOT EXISTS:

- `CREATE TABLE IF NOT EXISTS table_name(  
    column1 datatype,  
    column2 datatype,  
    column3 datatype,  
    .....  
    columnN datatype,  
    PRIMARY KEY( one or more columns )  
);`

- Creating a Table from an Existing Table:

- ```
CREATE TABLE NEW_TABLE_NAME AS
SELECT [column1, column2...columnN]
FROM EXISTING_TABLE_NAME
WHERE Condition;
```

- Example: Here NOT NULL ensures we do not insert null.

```
CREATE TABLE CUSTOMERS (
    ID INT NOT NULL,
    NAME VARCHAR (20) NOT NULL,
    AGE INT NOT NULL,
    ADDRESS CHAR (25) ,
    SALARY DECIMAL (18, 2),
    PRIMARY KEY (ID)
);
```

- **SQL INSERT INTO Statement:**

- ```
INSERT INTO table_name(
column1,column2....columnN)
VALUES ( value1, value2....valueN);
```
- Example: We can insert more rows at a time.

```
INSERT INTO CUSTOMERS VALUES
(1, 'Ramesh', 32, 'Ahmedabad', 2000.00 ),
(2, 'Khilan', 25, 'Delhi', 1500),
(3, 'kaushik', 23, 'Kota', 2000),
(4, 'Chaitali', 25, 'Mumbai', 6500),
(5, 'Hardik', 27, 'Bhopal', 8500),
(6, 'Komal', 22, 'Hyderabad', 4500),
(7, 'Muffy', 24, 'Indore', 10000);
```

- Inserting Data into a Table Using Another

- ```
INSERT INTO first_table_name
(column_name(s))
SELECT column1, column2, ...columnN
FROM second_table_name
[WHERE condition];
```
- ```
CREATE TABLE BUYERS (
    ID INT NOT NULL,
```

```

NAME VARCHAR (20) NOT NULL,
AGE INT NOT NULL,
ADDRESS CHAR (25) ,
SALARY DECIMAL (18, 2),
PRIMARY KEY (ID)
);

```

- ```
INSERT INTO BUYERS (ID, NAME, AGE, ADDRESS, SALARY)
SELECT * FROM CUSTOMERS;
```
- Insert whole Table:
- ```
INSERT INTO first_table_name TABLE
second_table_name;
```

○ **Insert Into... Select Statement:**

- ```
INSERT INTO table_new
SELECT (column1, column2, ...columnN)
FROM table_old;
```

○ **SQL SELECT Statement:**

- ```
SELECT column1, column2....columnN FROM
table_name;
```
- ```
SELECT * FROM CUSTOMERS; --> For Selecting All
Rows.
```

○ **SQL UPDATE Statement:**

- ```
UPDATE table_name SET column1 = value1, column2 =
value2....columnN=valueN WHERE CONDITION;
```

■ **Example:**

- ```
UPDATE CUSTOMERS SET ADDRESS = 'Pune' WHERE ID = 6;
```

■ **SQL RENAME TABLE Statement:**

- ```
RENAME TABLE table_name TO new_table_name;
```
- Or 

```
ALTER TABLE table_name RENAME [TO|AS]
new_table_name
```
- **Example:**
- ```
ALTER TABLE BUYERS RENAME TO CUSTOMERS;
```

○ **SQL DROP TABLE Statement:**

- **SQL DROP TABLE Is Used To Drop Tables.**
- ```
DROP TABLE table_name;
```

- SQL **DROP DATABASE** statement is used to delete an existing database along with all the data such as tables, views, indexes, stored procedures, and constraints.
- **DROP** is used to **completely delete a table**, including all of its data and structure.
- This will **completely delete** the **CUSTOMERS** table.
- You cannot access the table or any of its data after this.
  - `DROP DATABASE IF EXISTS DatabaseName;`
  - Deleting Multiple Databases
  - `DROP DATABASE testDB3, testDB4;`
- **SQL TRUNCATE TABLE Statement:**
  - The TRUNCATE TABLE statement is implemented in SQL to delete the data of the table but not the table itself.
  - When this SQL statement is used, the table stays in the database like an empty table.
  - `TRUNCATE TABLE table_name;`
- **DELETE – Remove Rows (Data)**
  - **DELETE** is used to remove specific rows (records) from a table, but keeps the table and its structure.
  - **DELETE FROM table\_name WHERE condition;**
  - **Example:**
    - **DELETE FROM CUSTOMERS WHERE CustomerID = 1;**
    - This will remove only the row where CustomerID = 1 from the CUSTOMERS table.
    - The table and other rows remain.
- **SQL ALTER TABLE Statement:**
  - `ALTER TABLE table_name`  
`{ADD|DROP|MODIFY} column_name {data_type};`
  - Example:
 

```
ALTER TABLE CUSTOMERS ADD SEX1 char(1);
```

```
ALTER TABLE table_name DROP COLUMN
column_name;
```
  - ALTER TABLE – RENAME COLUMN:
    - `ALTER TABLE table_name`  
`RENAME COLUMN old_column_name to`  
`new_column_name;`
  - `ALTER TABLE table_name RENAME TO new_table_name;`

## ○ SQL DISTINCT Clause:

- The DISTINCT clause in a database is used to identify the non-duplicate data from a column.
- ```
SELECT DISTINCT column1, column2....columnN FROM table_name;
```
- DISTINCT Keyword with COUNT() Function
  - ```
SELECT COUNT(DISTINCT column_name)
```

```
FROM table_name WHERE condition;
```

## ○ SQL WHERE Clause:

 Key Difference:

| Clause                | Filters...               | Used With                   | Happens...        |
|-----------------------|--------------------------|-----------------------------|-------------------|
| <code>WHERE</code>    | Individual rows          | Before grouping             | First             |
| ■ <code>HAVING</code> | Groups (aggregated data) | After <code>GROUP BY</code> | After aggregation |

- Note: We should use where before the group by.  

```
SELECT Customer, SUM(Quantity) AS TotalQuantity  
FROM Orders  
WHERE Product = 'Banana'  
GROUP BY Customer  
HAVING SUM(Quantity) >= 5;
```
- ```
SELECT column1, column2....columnN  
FROM table_name  
WHERE CONDITION;
```

## ○ SQL AND/OR Operators:

- ```
SELECT column1, column2....columnN  
FROM table_name  
WHERE CONDITION-1 {AND|OR} CONDITION-2;
```
- Example:
  - ```
SELECT ID, NAME, SALARY FROM CUSTOMERS WHERE  
SALARY > 2000 AND age < 25;
```

## ○ SQL IN/NOT IN Clause:

- ```
SELECT column1, column2....columnN  
FROM table_name  
WHERE column_name IN (val-1, val-2,...val-N);
```

- ```
SELECT column1, column2....columnN
FROM table_name
WHERE column_name NOT IN (val-1,
val-2,...val-N);
```

- Example:

- ```
SELECT * FROM CUSTOMERS
WHERE NAME IN ('Khilan', 'Hardik', 'Muffy');
```

## ○ SQL ANY/ALL Clause:

✓ Example:

Assume this table: `products`

price

200

300

500

sql

```
SELECT *
FROM products
WHERE price > ANY (SELECT price FROM products WHERE price < 300);
```

➡ This checks:

- `price > 200` → returns rows with price 300 and 500.

sql

```
SELECT *
FROM products
WHERE price > ALL (SELECT price FROM products WHERE price < 300);
```

- returns rows with price 300 and 500.

- ANY Clause Check if atleast Any one true - consider if we get 100,200 in above table price >any (100,200) means our price is 110 then 110 > any(100,200) → true as 110 >100.

- ALL Clause Check that all should pass. 110 > ALL (100,200) → false.

## ○ SQL BETWEEN Clause:

- ```
SELECT column1, column2....columnN
FROM table_name
```

```
WHERE column_name BETWEEN val-1 AND val-2;
```

- Example:

- ```
SELECT * FROM CUSTOMERS WHERE AGE BETWEEN 20 AND 25;
```

- **SQL LIKE Clause :**

- ```
SELECT column1, column2....columnN  
  
FROM table_name  
  
WHERE column_name LIKE PATTERN;
```

- We can also use NOT before LIKE.

- PATTERN - In this we use wildcards

- % — The percent sign represents zero, one or multiple characters. Ex:- 2% May be 20 or 200 or 220 or any other.
- \_ — The underscore represents a single number or character. Ex:- 2\_ May be 22 or 23 or 2a But not 222.

- Example:

- ```
SELECT * FROM CUSTOMERS WHERE SALARY LIKE '200%';
```

- **SQL ORDER BY Clause:**

- ```
SELECT column1, column2....columnN  
  
FROM table_name  
  
WHERE CONDITION  
  
ORDER BY column_name {ASC|DESC};
```

- Example:

- ```
SELECT * FROM CUSTOMERS ORDER BY NAME ASC;
```

- ORDER BY Clause on Multiple Columns:

- ```
SELECT * FROM CUSTOMERS ORDER BY AGE ASC,  
SALARY DESC;
```

- ORDER BY with LIMIT Clause

- ```
SELECT column1, column2, ...
```

```
FROM table_name
```

```
ORDER BY column_name1 [ASC | DESC],  
column_name2 [ASC | DESC], ...
```

```
LIMIT N;
```



- ORDER BY with WHERE Clause

- `SELECT * FROM CUSTOMERS`

- `WHERE AGE = 25 ORDER BY NAME DESC;`

- We cannot use **ASC** or **DESC** without **ORDER BY** in SQL.

- **SQL COUNT Function:**

- The COUNT Function gives the number of non-null values present in the specified column.

- `SELECT COUNT(column_name)`

- `FROM table_name;`

- **SQL GROUP BY Clause:**

- **GROUP BY** is used in SQL to group rows that have the same values in one or more columns, so that we can apply aggregate functions (like **COUNT**, **SUM**, **AVG**, etc.) to each group.
  - `SELECT column1, AGGREGATE_FUNCTION(column2)`

`FROM table_name`

`GROUP BY column1;`

**Example:**

 **Example Table: Sales**

| ID | Product | Quantity |
|----|---------|----------|
| 1  | Apple   | 10       |
| 2  | Banana  | 5        |
| 3  | Apple   | 7        |
| 4  | Banana  | 3        |
| 5  | Orange  | 8        |

- `SELECT Product, SUM(Quantity)`

`FROM Sales`

`GROUP BY Product;`

### Output:

| Product | SUM(Quantity) |
|---------|---------------|
| Apple   | 17            |
| Banana  | 8             |
| Orange  | 8             |

○

○ **Example: For Grouping Two columns.**

| ID | Product | Region | Quantity |
|----|---------|--------|----------|
| 1  | Apple   | East   | 10       |
| 2  | Banana  | West   | 5        |
| 3  | Apple   | East   | 7        |
| 4  | Banana  | East   | 3        |
| 5  | Apple   | West   | 8        |

○

6 Banana West 6

○ **SELECT Product, Region, SUM(Quantity) AS TotalQty**

**FROM Sales**

**GROUP BY Product, Region;**

### Output:

| Product | Region | TotalQty |
|---------|--------|----------|
| Apple   | East   | 17       |
| Apple   | West   | 8        |
| Banana  | East   | 3        |
| Banana  | West   | 11       |

○

○ **GROUP BY with ORDER BY Clause**

```
■ SELECT column1, column2, ...,  
    aggregate_function(columnX) AS alias  
FROM table  
GROUP BY column1, column2, ...
```

```
ORDER BY column1 [ASC | DESC], column2 [ASC  
| DESC], ...;
```

■ Example:

```
■ SELECT AGE, MIN(SALARY) AS MIN_SALARY  
  
FROM CUSTOMERS  
  
GROUP BY AGE ORDER BY MIN_SALARY DESC;
```

● SQL HAVING Clause

- What is HAVING?
- The HAVING clause is used to filter groups created by the GROUP BY clause.
- It allows you to apply conditions on aggregate functions like SUM(), COUNT(), AVG(), MIN(), MAX() etc.
- Think of HAVING as a WHERE clause for groups.

```
SELECT column1, column2, ..  
  
aggregate_function(column_Name)  
  
FROM table_name  
  
GROUP BY column1, column2 ..  
  
HAVING condition;
```

Important Notes:

- HAVING is applied after aggregation.
- If you want to filter rows before aggregation, use WHERE.
- You can use HAVING without GROUP BY – it will treat the entire result as a single group
- Note: We should use having after the group by.  
SELECT Customer, SUM(Quantity) AS TotalQuantity  
FROM Orders  
WHERE Product = 'Banana'  
GROUP BY Customer  
HAVING SUM(Quantity) >= 5;
- The following code block shows the position of the HAVING Clause - when we write a query first select should come after that from and where and so on.
  1. SELECT
  2. FROM
  3. Joins

4. WHERE
5. GROUP BY
6. HAVING
7. ORDER BY

- **SQL Constraints:**

- SQL Constraints are the rules applied to a data columns or the complete table to limit the type of data that can go into a table.
- **SQL Create Constraints**

- ```
CREATE TABLE table_name (
    column1 datatype constraint,
    column2 datatype constraint,
    ....
    columnN datatype constraint
);
```

- **Example**

1. 

```
CREATE TABLE CUSTOMERS (
    ID INT NOT NULL UNIQUE,
    NAME VARCHAR (20) NOT NULL DEFAULT 'Not Available',
    AGE INT NOT NULL CHECK (AGE>=18),
    ADDRESS CHAR (25),
    SALARY DECIMAL (18, 2),
    PRIMARY KEY (ID),
    CUSTOMER INT FOREIGN KEY REFERENCES CUSTOMER_ID (ID)
);
```

2. **Here In** Above NOT NULL, UNIQUE, DEFAULT, CHECK, PRIMARY KEY, FOREIGN KEY ARE **Constraints.**

- **What is an SQL View?**

- An SQL View is a virtual table based on the result of a **SELECT** query. It does not store data physically; instead, it displays data dynamically from one or more real tables.

- **CREATE VIEW view\_name AS**

- SELECT column1, column2....**

- FROM table\_name**

- WHERE [condition];**

- Update View Table:

- **UPDATE view\_name**

- SET column1 = value1, column2 = value2...,**  
**columnN = valueN**

- WHERE [condition];**

- Drop view table:[] → optional

- **DROP VIEW [IF EXISTS] view\_name;**

- **DELETE, RENAME** → Same like a normal table.

- **SQL TOP Clause:**

- **SELECT TOP value column\_name(s)**

- FROM table\_name**

- WHERE [condition];**

- Example: Return top 4 and 40 percent rows.

- **SELECT TOP 4 \* FROM CUSTOMERS;**

- **SELECT TOP 40 PERCENT \* FROM CUSTOMERS ORDER**  
**BY SALARY;**

- **DELETE TOP (2) FROM CUSTOMERS WHERE NAME LIKE**  
**'K%';**

- **SQL EXISTS** Clause:

Customers :

| customer_id | name    |
|-------------|---------|
| 1           | Alice   |
| 2           | Bob     |
| 3           | Charlie |

Orders :

| order_id | customer_id | product |
|----------|-------------|---------|
| 101      | 1           | Phone   |
| 102      | 2           | Laptop  |

○

○ **SELECT name**

**FROM Customers c**

**WHERE EXISTS (**

**SELECT 1**

**FROM Orders o**

**WHERE o.customer\_id = c.customer\_id**

**);--here we can select any thing (we choice 1)**

○

**Here we select a row if the inner query returns true.we move from first to last row for every row we check the inner query.**

✅ **Final Output:**

| name  |
|-------|
| Alice |
| Bob   |

○

● **SQL CASE Clause:**

○

**CASE**

**WHEN condition1 THEN statement1,**

**WHEN condition2 THEN statement2,**

**WHEN condition THEN statementN**

**ELSE result -- optional by default id return null**

**END ;**

| ID | NAME     | AGE | ADDRESS   | SALARY   |
|----|----------|-----|-----------|----------|
| 1  | Ramesh   | 32  | Ahmedabad | 2000.00  |
| 2  | Khilan   | 25  | Delhi     | 1500.00  |
| 3  | Kaushik  | 23  | Kota      | 2000.00  |
| 4  | Chaitali | 25  | Mumbai    | 6500.00  |
| 5  | Hardik   | 27  | Bhopal    | 8500.00  |
| 6  | Komal    | 22  | Hyderabad | 4500.00  |
| 7  | Muffy    | 24  | Indore    | 10000.00 |

○ **SELECT NAME, AGE,**

**CASE**

**WHEN AGE > 30 THEN 'Gen X'**

**WHEN AGE > 25 THEN 'Gen Y'**

**WHEN AGE > 22 THEN 'Gen Z'**

**ELSE 'Gen Alpha'**

**END AS Generation**

**FROM CUSTOMERS;**

**Output:-**

| NAME     | AGE | Generation |
|----------|-----|------------|
| Ramesh   | 32  | Gen X      |
| Khilan   | 25  | Gen Z      |
| Kaushik  | 23  | Gen Z      |
| Chaitali | 25  | Gen Z      |
| Hardik   | 27  | Gen Y      |
| Komal    | 22  | Gen Alpha  |
| Muffy    | 24  | Gen Z      |

- **SQL UNION/UNION ALL Operator**

- To use the UNION operator on multiple tables, all these tables must be union compatible.
- Both used to retrieve the rows from multiple tables and return them as one single table.

- The difference between these two operators is that UNION only returns distinct rows while UNION ALL returns all the rows present in the tables.
  - The same number of columns selected with the same datatype.
  - These columns must also be in the same order.
  - They need not have the same number of rows.
- *The column names in the final result set will be based on the column names selected in the first SELECT statement. If you want to use a different name for a column in the final result set, you can use an alias in the SELECT statement.*

- `SELECT column1 [, column2 ]`

`FROM table1 [, table2 ]`

`[WHERE condition]`

`UNION`

`SELECT column1 [, column2 ]`

`FROM table1 [, table2 ]`

`[WHERE condition];`

- *Union All:*

`SELECT * FROM table1`

`UNION ALL`

`SELECT * FROM table2;`

- *Union with order by:*

`SELECT ID, SALARY FROM CUSTOMERS WHERE ID > 5`

`UNION`

`SELECT CUSTOMER_ID, AMOUNT FROM ORDERS WHERE  
CUSTOMER_ID > 2`

`ORDER BY SALARY;`

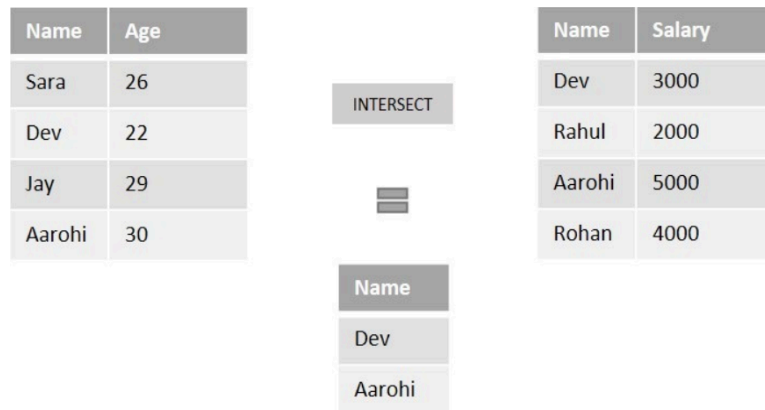
- *Example:*

- `SELECT SALARY FROM CUSTOMERS UNION SELECT  
AMOUNT FROM ORDERS;`

- **SQL INTERSECT Operator:**



- The **INTERSECT** operator in SQL is used to retrieve the records that are identical/common between the result sets of two or more tables.



- 

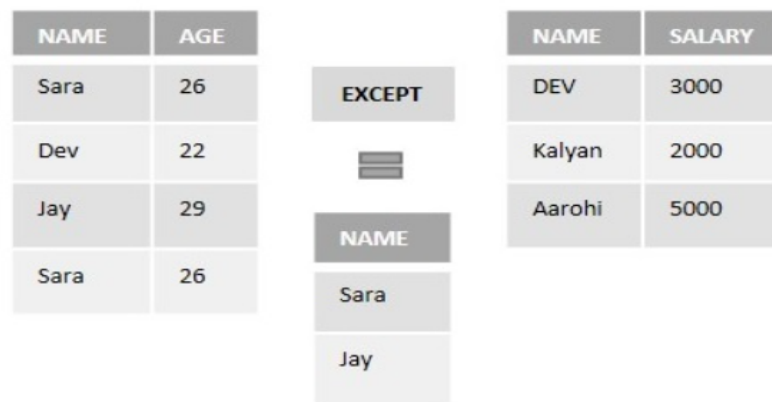
```
SELECT column1, column2,..., columnN
FROM table1
INTERSECT
SELECT column1, column2,..., columnN
FROM table2
```

- Example:

```
■ SELECT NAME, AGE, HOBBY FROM STUDENTS HOBBY
INTERSECT
SELECT NAME, AGE, HOBBY FROM STUDENTS;
```

## • The SQL EXCEPT Operator

- The **EXCEPT** operator in SQL is used to retrieve all the unique records from the left operand (query), except the records that are present in the result set of the right operand (query).



- 

```
SELECT column1, column2,..., columnN
```

```
FROM table1, table2,..., tableN
```

```
[Conditions] //optional
```

```
EXCEPT
```

```
SELECT column1, column2,..., columnN
```

```
FROM table1, table2,..., tableN
```

```
[Conditions] //optional
```

Example:

```
SELECT NAME, HOBBY, AGE FROM STUDENTS
```

```
EXCEPT
```

```
SELECT NAME, HOBBY, AGE FROM STUDENTS_HOBBY;
```

- **What are SQL Data types?**

- An SQL data type refers to the type of data which can be stored in a column of a database table.
- For example, if you want to store student name in a column then you should give column name something like student\_name and its data type will be char(50) which means it can store a string of characters up to 50 characters.
- The data type provides guidelines for SQL to understand what type of data is expected inside each column, and hence, prevents the user from entering any unexpected or invalid data in a column.
- Types of SQL Data Types:
  - Character/String Data Types:
    1. CHAR(n) → If we store a data whose size is less than n, then it is padded.  $0 \leq n \leq 255$ 
      - Example:
      - CREATE TABLE ExampleChar (Code CHAR(5));
      - INSERT INTO ExampleChar VALUES ('USA');
      - Stored as 'USA ' (padded with 2 spaces to make length 5).
      - Length will always be 5, even if data is shorter.
    2. VARCHAR(n) → Here no padding is there.  $0 \leq n \leq 65535$ 
      - CREATE TABLE ExampleVarchar (Name VARCHAR(5));
      - INSERT INTO ExampleVarchar VALUES ('USA');
      - Stored as 'USA' (only 3 bytes, no padding).

- Length depends on actual string length.

## ■ Numeric Data Types:

1. **INT** → A normal-sized integer that can be signed or unsigned. If signed, the allowable range is from -2147483648 to 2147483647. If unsigned, the allowable range is from 0 to 4294967295.
  - `CREATE TABLE Students (ID INT, Age INT);`
  - → Example: ID = 1, Age = 20
2. **DECIMAL(p, s)** → Fixed-precision numbers.
  - `CREATE TABLE Products ( Price DECIMAL(6,2));`
  - → Example: Price = 1234.56 (6 digits total, 2 after decimal)
3. **FLOAT** → Approximate decimal values.
  - `CREATE TABLE Physics ( Gravity FLOAT);`
  - → Example: Gravity = 9.81
4. **TEXT** → Used for long text.
  - `CREATE TABLE Articles (Content TEXT);`
  - → Example: 'This is a full article body with many characters...'

## ■ Date and Time Types:

1. **DATE** → Stores calendar date (YYYY-MM-DD)
  - `CREATE TABLE Employees (JoinDate DATE);`
  - → Example: '2023-01-15'
2. **TIME** → Stores time of day (HH:MM:SS)
  - `CREATE TABLE Schedules (StartTime TIME);`
  - → Example: '08:30:00'
3. **DATETIME or TIMESTAMP** → Stores both date and time.
4. `CREATE TABLE Events (EventTime DATETIME);`
5. → Example: '2025-05-23 14:00:00'

## ■ Boolean Type:

1. **BOOLEAN or BOOL**
  - Stores TRUE or FALSE values.
  - `CREATE TABLE Members (IsActive BOOLEAN);`
  - Example: IsActive = TRUE

## ■ What is **ENUM**: **ENUM** is a **string data type** that allows you to **define a list of allowed values** for a column.

1. It ensures that the column can **only store one of those values** — like a built-in dropdown list!
2. `CREATE TABLE Shirts (Size ENUM('S', 'M', 'L', 'XL'));`

- The **Size** column can only contain one of these: 'S', 'M', 'L', or 'XL'.

## • What is SQL Operator?

- Types of Operator in SQL:

### ■ SQL Arithmetic Operators:

|   |                |               |
|---|----------------|---------------|
| + | Addition       | 10 + 20 = 30  |
| - | Subtraction    | 20 - 30 = -10 |
| * | Multiplication | 10 * 20 = 200 |
| / | Division       | 20 / 10 = 2   |
| % | Modulus        | 5 % 2 = 1     |

1.

### ■ SQL Comparison Operators:

|    |                          |                      |
|----|--------------------------|----------------------|
| =  | Equal to                 | 5 = 5 returns TRUE   |
| != | Not equal                | 5 != 6 returns TRUE  |
| <> | Not equal                | 5 <> 4 returns TRUE  |
| >  | Greater than             | 4 > 5 returns FALSE  |
| <  | Less than                | 4 < 5 returns TRUE   |
| >= | Greater than or equal to | 4 >= 5 returns FALSE |
| <= | Less than or equal to    | 4 <= 5 returns TRUE  |
| !< | Not less than            | 4 !< 5 returns FALSE |
| !> | Not greater than         | 4 !> 5 returns TRUE  |

1.

### ■ SQL Logical Operators:

|         |                                                                                             |
|---------|---------------------------------------------------------------------------------------------|
| ALL     | TRUE if all of a set of comparisons are TRUE.                                               |
| AND     | TRUE if all the conditions separated by AND are TRUE.                                       |
| ANY     | TRUE if any one of a set of comparisons are TRUE.                                           |
| BETWEEN | TRUE if the operand lies within the range of comparisons.                                   |
| EXISTS  | TRUE if the subquery returns one or more records                                            |
| IN      | TRUE if the operand is equal to one of a list of expressions.                               |
| LIKE    | TRUE if the operand matches a pattern specially with wildcard.                              |
| NOT     | Reverses the value of any other Boolean operator.                                           |
| OR      | TRUE if any of the conditions separated by OR is TRUE                                       |
| IS NULL | TRUE if the expression value is NULL.                                                       |
| SOME    | TRUE if some of a set of comparisons are TRUE.                                              |
| UNIQUE  | The UNIQUE operator searches every row of a specified table for uniqueness (no duplicates). |

1.

## • SQL - Comments

- There are two types of comments used.

### ■ Single-line comments

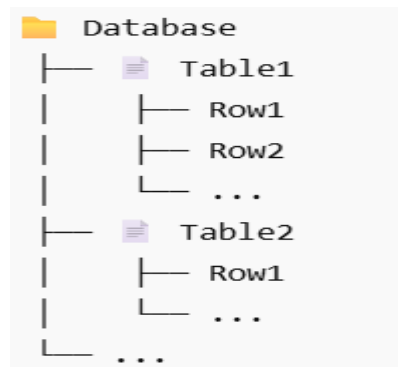
- `-- Will fetch all the table records`

```
SELECT * from table;
```

- Multi-line comments

1. `/* This is a`  
`multi-line`  
`comment */`

- DataBase Vs Tables



- **What is an Index?**

- An **index** is a **database structure** that improves the **speed of data retrieval** on a table.
  - Without an index: Full table scan
  - With an index: Jump directly to needed rows
- **How Does an Index Work (Internally)?**
  - Most SQL databases use B-Trees (or B+ Trees) (B → Balanced) for indexes.
  - **Benefits:**
    - **Logarithmic time** complexity:  $O(\log n)$
    - **Efficient for range queries**
    - Great for large datasets (even millions of rows)
  - `CREATE INDEX idx_name ON Employees(Name);`
  - Now this query is very fast: (Here Employees → Table\_Name and Name Column\_Name)
  - `SELECT * FROM Employees WHERE Name = 'Alice';`
  - `DROP INDEX index_name ON table_name;`
  - We can declare Columns as indexes → Composite Index.
  - **Types:**



## 4. Types of Indexes

| Type                | Description                                               |
|---------------------|-----------------------------------------------------------|
| Primary Index       | Created automatically on primary key                      |
| Unique Index        | Ensures all values are unique                             |
| Composite Index     | Index on multiple columns                                 |
| Clustered Index     | Sorts the actual table rows (1 per table in SQL Server)   |
| Non-clustered Index | A separate structure with pointers to rows                |
| Full-text Index     | For searching text (like documents)                       |
| Bitmap Index        | Efficient for low-cardinality columns (used in analytics) |
| Spatial Index       | For geographic data (used in GIS systems)                 |

### ■ When Should You Use Indexes?

#### 1. ✓ Use Indexes:

- On columns used in **WHERE, JOIN, ORDER BY, GROUP BY**
- On foreign key columns
- On frequently searched or sorted fields

#### • ✗ Avoid Indexes:

- On columns with many duplicate values (e.g. gender)
- On columns that are rarely queried
- On very small tables — scanning is fast anyway
- Too many indexes can slow **INSERT/UPDATE/DELETE** because they must also update the index

### • SQL Join Clause:

- A **JOIN** in SQL is used to **combine rows** from two or more tables, based on a **related column** between them.
- (OR) SQL **Join** clause is used to combine data from two or more tables in a database. When the related data is stored across multiple tables, joins help you to retrieve records combining the fields from these tables using their foreign keys.
- *When you retrieve a table using joins, the resultant table displayed is not stored anywhere in the database.*
- JOINS --- (I) INNER JOIN (II) OUTER JOIN – (a) LEFT JOIN (b) RIGHT JOIN (c) FULL JOIN (III) OTHER – (a) SELF JOIN (b) CROSS JOIN.
- Let Consider An Example For Better Understand:
  - Table 1:

● Table: Customers

| customer_id | name  |
|-------------|-------|
| 1           | Alice |
| 2           | Bob   |
| 3           | Carol |
| 4           | David |

○ Table 2:

● Table: Orders

| order_id | customer_id | product  |
|----------|-------------|----------|
| 101      | 1           | Phone    |
| 102      | 2           | Laptop   |
| 103      | 2           | Keyboard |
| 104      | 5           | Tablet   |

○ ① INNER JOIN

■ Code :

- SELECT \*  
FROM Customers  
INNER JOIN Orders  
ON Customers.customer\_id =  
Orders.customer\_id;

● Result:

| customer_id | name  | order_id | customer_id (Orders) | product  |
|-------------|-------|----------|----------------------|----------|
| 1           | Alice | 101      | 1                    | Phone    |
| 2           | Bob   | 102      | 2                    | Laptop   |
| 2           | Bob   | 103      | 2                    | Keyboard |

● ② LEFT JOIN

○ Code:

- `SELECT *`  
`FROM Customers`  
`LEFT JOIN Orders`  
`ON Customers.customer_id =`  
`Orders.customer_id;`

● Result:

| customer_id | name  | order_id | customer_id (Orders) | product  |
|-------------|-------|----------|----------------------|----------|
| 1           | Alice | 101      | 1                    | Phone    |
| 2           | Bob   | 102      | 2                    | Laptop   |
| 2           | Bob   | 103      | 2                    | Keyboard |
| 3           | Carol | NULL     | NULL                 | NULL     |
| 4           | David | NULL     | NULL                 | NULL     |

## • **3 RIGHT JOIN**

- **Code**

- `SELECT *`  
`FROM Customers`  
`RIGHT JOIN Orders`  
`ON Customers.customer_id =`  
`Orders.customer_id;`

● Result:

| customer_id (Customers) | name  | order_id | customer_id | product  |
|-------------------------|-------|----------|-------------|----------|
| 1                       | Alice | 101      | 1           | Phone    |
| 2                       | Bob   | 102      | 2           | Laptop   |
| 2                       | Bob   | 103      | 2           | Keyboard |
| NULL                    | NULL  | 104      | 5           | Tablet   |

## • **4 FULL OUTER JOIN (Simulated with UNION)**

- **Code**

- `SELECT *`  
`FROM Customers`  
`LEFT JOIN Orders ON`  
`Customers.customer_id =`  
`Orders.customer_id`



```

UNION
SELECT *
FROM Customers
RIGHT JOIN Orders ON
Customers.customer_id =
Orders.customer_id;

```

● Result:

| customer_id | name  | order_id | customer_id (Orders) | product  |
|-------------|-------|----------|----------------------|----------|
| 1           | Alice | 101      | 1                    | Phone    |
| 2           | Bob   | 102      | 2                    | Laptop   |
| 2           | Bob   | 103      | 2                    | Keyboard |
| 3           | Carol | NULL     | NULL                 | NULL     |
| 4           | David | NULL     | NULL                 | NULL     |
| NULL        | NULL  | 104      | 5                    | Tablet   |

## • 5 CROSS JOIN

### ○ Code

```

■ SELECT *
FROM Customers
CROSS JOIN Orders;

```

● Result: (only first 5 rows shown here)

| customer_id | name  | order_id | customer_id (Orders) | product  |
|-------------|-------|----------|----------------------|----------|
| 1           | Alice | 101      | 1                    | Phone    |
| 1           | Alice | 102      | 2                    | Laptop   |
| 1           | Alice | 103      | 2                    | Keyboard |
| 1           | Alice | 104      | 5                    | Tablet   |
| 2           | Bob   | 101      | 1                    | Phone    |
| ...         | ...   | ...      | ...                  | ...      |

○ | Total rows = 4 customers × 4 orders = 16 rows

## • SQL DELETE... JOIN Clause:

- The purpose of Joins in SQL is to combine records of two or more tables based on common columns/fields. Once the tables are joined,

performing the deletion operation on the obtained result-set will delete records from all the original tables at a time.

- For example, consider a database of an educational institution. It consists of various tables: Departments, StudentDetails, LibraryPasses, LaboratoryPasses etc. When a set of students are graduated, all their details from the organizational tables need to be removed, as they are unwanted. However, removing the details separately from multiple tables can be cumbersome.
- Syntax:

- ```
DELETE table(s)
FROM table1 JOIN table2
ON table1.common_field =
table2.common_field;
```

- Example:

- ```
DELETE a
FROM CUSTOMERS AS a INNER JOIN ORDERS AS b
ON a.ID = b.CUSTOMER_ID;
```